

Name: Dipak Shingade Rollno.: 21102B0029 BE CMPN B

---

## Experiment 1 : Predict Housing Prices using Linear Regression

---

Github : <https://github.com/Dipakgith/ML-lab-experiments>

---

Dataset: <https://www.kaggle.com/datasets/camnugent/california-housing-prices>

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import Pipeline
from sklearn.metrics import r2_score, mean_absolute_error,
mean_squared_error

import warnings
warnings.filterwarnings('ignore')

housing = pd.read_csv(r"/content/housing.csv")
```

### Data Preprocessing

```
print(housing.head(10))
```

	longitude	latitude	housing	median age	total rooms
total bedrooms \					
0	-122.23	37.88		41.0	880.0
129.0					
1	-122.22	37.86		21.0	7099.0
1106.0					
2	-122.24	37.85		52.0	1467.0
190.0					
3	-122.25	37.85		52.0	1274.0
235.0					
4	-122.25	37.85		52.0	1627.0
280.0					
5	-122.25	37.85		52.0	919.0
213.0					
6	-122.25	37.84		52.0	2535.0

```

489.0
7    -122.25    37.84    52.0    3104.0
687.0
8    -122.26    37.84    42.0    2555.0
665.0
9    -122.25    37.84    52.0    3549.0
707.0

```

	population	households	median income	median house value
ocean_proximity				
0	322.0	126.0	8.3252	452600.0
NEAR BAY				
1	2401.0	1138.0	8.3014	358500.0
NEAR BAY				
2	496.0	177.0	7.2574	352100.0
NEAR BAY				
3	558.0	219.0	5.6431	341300.0
NEAR BAY				
4	565.0	259.0	3.8462	342200.0
NEAR BAY				
5	413.0	193.0	4.0368	269700.0
NEAR BAY				
6	1094.0	514.0	3.6591	299200.0
NEAR BAY				
7	1157.0	647.0	3.1200	241400.0
NEAR BAY				
8	1206.0	595.0	2.0804	226700.0
NEAR BAY				
9	1551.0	714.0	3.6912	261100.0
NEAR BAY				

```
housing.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
-----
0   longitude             20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age    20640 non-null  float64
3   total_rooms           20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population            20640 non-null  float64
6   households            20640 non-null  float64
7   median_income         20640 non-null  float64
8   median_house_value    20640 non-null  float64
9   ocean_proximity       20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

```
print(housing.describe())
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

```
print(housing.nunique())
```

longitude	844
latitude	862
housing_median_age	52
total_rooms	5926
total_bedrooms	1923
population	3888
households	1815
median_income	12928
median_house_value	3842
ocean_proximity	5
dtype:	int64

```
print(housing.isnull().sum())
```

longitude	0
latitude	0
housing_median_age	0

```

total_rooms      0
total_bedrooms   207
population        0
households        0
median_income     0
median_house_value 0
ocean_proximity  0
dtype: int64

housing['total_bedrooms'].fillna(method='ffill', inplace=True)
print(housing.isnull().sum())
print(housing.duplicated().sum())

longitude        0
latitude          0
housing_median_age 0
total_rooms       0
total_bedrooms    0
population        0
households        0
median_income     0
median_house_value 0
ocean_proximity  0
dtype: int64
0


```

## Visualization

```

housing.plot(kind='box', subplots=True, layout=(3,3), figsize=(15,7))
plt.show()

```

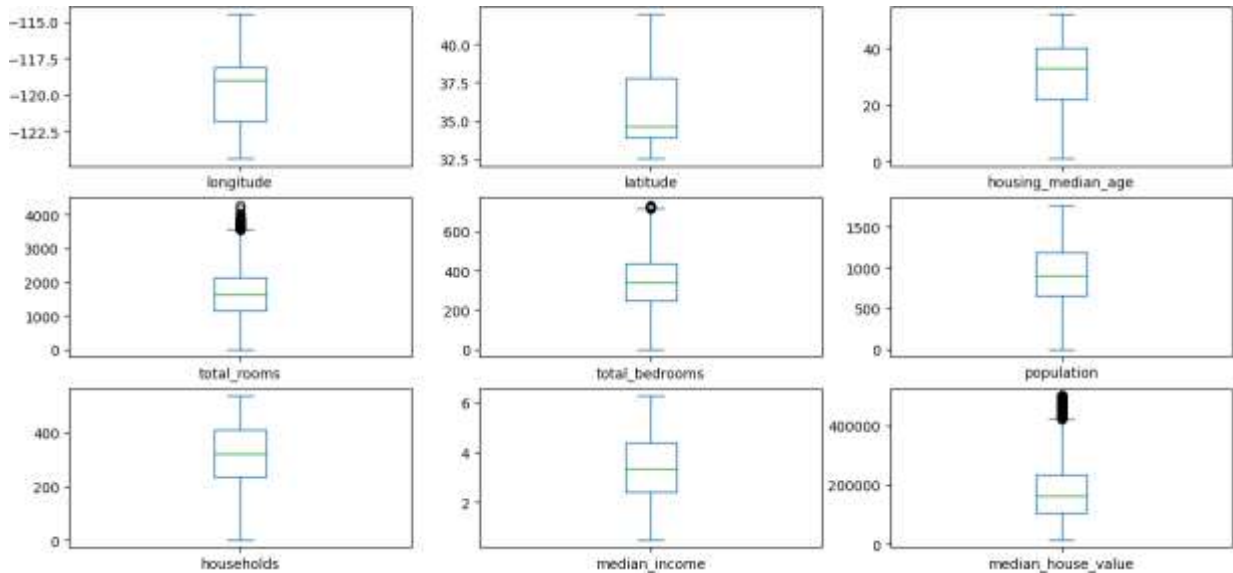


```

housing1 = housing.copy()
for col in ['total_rooms', 'total_bedrooms', 'population',
            'households', 'median_income']:
    housing1 = housing1[housing1[col] < housing1[col].quantile(0.9)]

housing1.plot(kind='box', subplots=True, layout=(3,3), figsize=(15,7))
plt.show()


```



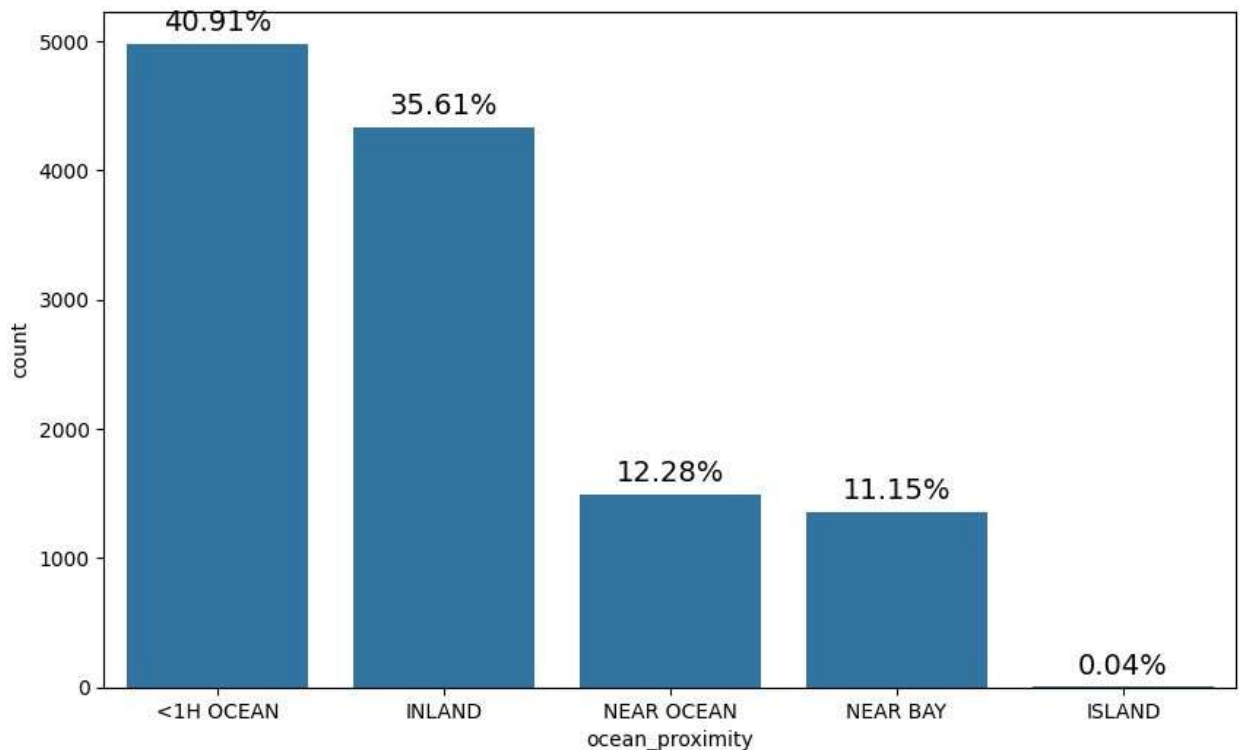
```

housing1.hist(bins=50, figsize=(20,15))
plt.show()

```



```
ocean_values = housing1["ocean_proximity"].value_counts()
plt.figure(figsize=(10,6))
sns.countplot(x="ocean_proximity", data=housing1,
order=ocean_values.index)
for i in range(ocean_values.shape[0]):
    count = ocean_values[i]
    strt = '{:0.2f}%'.format(100 * count / housing1.shape[0])
    plt.text(i, count + 100, strt, ha='center', color='black',
    fontsize=14)
plt.show()
```



```
import urllib.request
import io
import matplotlib.image as mpimg

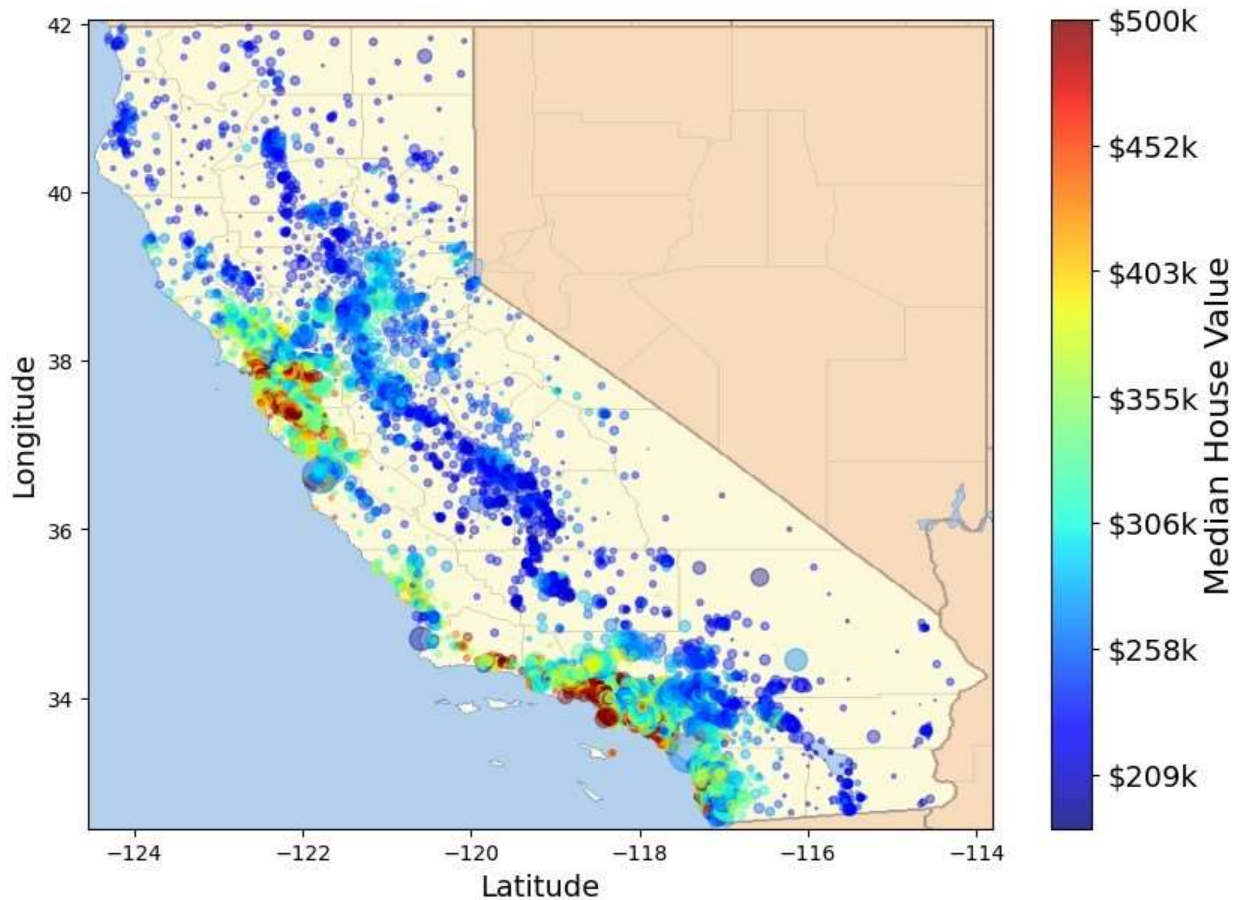
DOWNLOAD_ROOT = "https://raw.githubusercontent.com/ageron/handson-ml2/master/"
filename = "california.png"
print("Downloading", filename)
url = DOWNLOAD_ROOT + "images/end_to_end_project/" + filename
with urllib.request.urlopen(url) as url_request:
    image_data = url_request.read()

image_data = io.BytesIO(image_data)
california_img = mpimg.imread(image_data, format='png')

ax = housing.plot(kind='scatter', x='longitude', y='latitude',
figsize=(10,7), s=housing['population']/100,
                    c='median_house_value', colorbar=False,
cmap=plt.get_cmap('jet'), alpha=0.4)
plt.imshow(california_img, alpha=0.8, extent=[-124.55, -113.80, 32.45,
42.05], cmap=plt.get_cmap('jet'))
plt.xlabel('Latitude', fontsize=14)
plt.ylabel('Longitude', fontsize=14)
prices = housing["median_house_value"]
tick_values = np.linspace(prices.min(), prices.max(), 11)
cbar = plt.colorbar(ticks=tick_values/prices.max())
```

```
cbar.ax.set_yticklabels(["$%dk" % (round(v/1000)) for v in
tick_values], fontsize=14)
cbar.set_label('Median House Value', fontsize=16)
plt.show()
```

Downloading california.png



## Model Generation

```
df = housing[['housing_median_age', 'total_rooms', 'total_bedrooms',
'population', 'households', 'median_income', 'median_house_value']]
X = df[['housing_median_age', 'total_rooms', 'total_bedrooms',
'population', 'households', 'median_income']]
y = df['median_house_value']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('regressor', LinearRegression())
])
pipeline.fit(X_train, y_train)
```



```
Pipeline(steps=[('scaler', StandardScaler()),
                  ('regressor', LinearRegression())])

scores = cross_val_score(pipeline, X_train, y_train,
                          scoring='neg_mean_squared_error', cv=10)
print('Cross-validation RMSE:', np.sqrt(-scores).mean())

Cross-validation RMSE: 75917.02514172823

print('Train score:', pipeline.score(X_train, y_train))
print('Test score:', pipeline.score(X_test, y_test))

Train score: 0.5700818543102877
Test score: 0.5420853869482856

y_pred = pipeline.predict(X_test)

print('R2 score:', r2_score(y_test, y_pred))
print('Mean Absolute Error:', mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', mean_squared_error(y_test, y_pred))

R2 score: 0.5420853869482856
Mean Absolute Error: 56836.05528278542
Mean Squared Error: 6000554202.106741
```